

Google Cloud deployment runbook \Docker Compose app on a small VM\)

This runbook captures the **prerequisites** and the **exact commands** to deploy a multi-container Docker Compose application (UI + API + DB) to Google Cloud using a low-cost Compute Engine VM (fits well within trial credits).

It is written for **Windows PowerShell**.

Prerequisites

1) Accounts \& billing

- A Google Cloud account with **Billing enabled** (trial credits are OK).
- A Google Cloud **Project ID** (example used below: `iron-tea-482716-v0`).

2) Local tools \Windows\)

Install Git

- Install Git for Windows and ensure `git` is available in PowerShell.

Verify:

```
git --version
```

Install Docker Desktop (optional for local run)

- Install Docker Desktop if you want to run locally. (Not required to deploy to GCP VM.)

Install Google Cloud CLI

Install Google Cloud CLI.

After installation, verify gcloud is callable.

On some Windows setups, the easiest is to call it via full path:

```
$gcloud = "$env:LOCALAPPDATA\Google\Cloud SDK\google-cloud-sdk\bin\gcloud.cmd"
& $gcloud --version
```

3) Your application source code \Git repo\)

You need a Git repo URL for your app, for example:

- `https://github.com/leenaparik/shivam-demo.git`

One\time setup \per machine / per account

Login to Google Cloud

```
$gcloud = "$env:LOCALAPPDATA\Google\Cloud SDK\google-cloud-sdk\bin\gcloud.cmd"  
& $gcloud auth login  
& $gcloud auth list
```

\Optional\ List projects you have access to

```
& $gcloud projects list --format="table(projectId,name)" --limit=50
```

Deployment steps \the commands

0) Choose your project + zone

```
$PROJECT_ID = "iron-tea-482716-v0"  
$ZONE = "us-central1-a"
```

1) Configure gcloud to use the project/zone

```
& $gcloud --quiet config set project $PROJECT_ID  
& $gcloud config set compute/zone $ZONE  
& $gcloud config list
```

2) Enable Compute Engine API

```
& $gcloud --quiet services enable compute.googleapis.com
```

3) Create a firewall rule for your app ports

Recommended: restrict to your IP (safer).

If you *must* open to all, use 0.0.0.0/0.

Open ports:

- UI: `8080`
- API: `5000`

```
$FW_RULE = "shivam-demo-allow-web"  
$TAGS = "shivam-demo"  
$SOURCE_RANGES = "0.0.0.0/0"  
& $gcloud --quiet compute firewall-rules create $FW_RULE `
```

```
--direction=INGRESS `  
--priority=1000 `  
--network=default `  
--action=ALLOW `  
--rules=tcp:8080,tcp:5000 `  
--source-ranges=$SOURCE_RANGES `  
--target-tags=$TAGS
```

If the rule already exists and you want to update allowed ports:

```
& $gcloud --quiet compute firewall-rules update $FW_RULE --rules=tcp:8080,tcp:5000
```

4) Create a small VM (e2-micro) and deploy via startup script

This approach uses a **startup script** to:

- install Docker + docker-compose
- clone your Git repo
- create `.env` with generated keys
- run `docker-compose up -d --build`

In this repo, the startup script file is:

- `gcp-startup.sh`

Create the VM:

```
$VM = "shivam-demo-vm"  
  
& $gcloud --quiet compute instances create $VM `  
  --machine-type=e2-micro `  
  --image-family=ubuntu-2204-lts `  
  --image-project=ubuntu-os-cloud `  
  --boot-disk-size=20GB `  
  --tags=$TAGS `  
  --metadata-from-file=startup-script=gcp-startup.sh
```

5) Get the VM external IP

```
$IP = & $gcloud compute instances describe $VM --zone $ZONE --format="value(networkInterfaces[0].accessConfig.0natIP)"  
$IP
```

Open the app:

- UI: `http://\$IP:8080`
- API health: `http://\$IP:5000/api/health`

6) Verify endpoints from PowerShell

```
curl.exe -s "http://$IP:5000/api/health"  
curl.exe -s -I "http://$IP:8080/" | findstr /I "HTTP"
```

Updating an existing deployment

If you changed code in GitHub and want the VM to redeploy, update the startup script metadata and reboot:

```
& $gcloud --quiet compute instances add-metadata $VM --zone $ZONE --metadata-from-file startup-script=gcp-startup.sh  
& $gcloud --quiet compute instances reset $VM --zone $ZONE
```

Basic troubleshooting commands

Check VM status

```
& $gcloud compute instances describe $VM --zone $ZONE --format="value(status)"
```

Read startup script logs (serial port output)

```
& $gcloud compute instances get-serial-port-output $VM --zone $ZONE --start=0
```

SSH to the VM and check containers

```
& $gcloud compute ssh $VM --zone $ZONE --command "sudo docker ps"  
& $gcloud compute ssh $VM --zone $ZONE --command "sudo docker ps -a"
```

Cost control (stay within free/trial credit)

Stop the VM when not using it

```
& $gcloud --quiet compute instances stop $VM --zone $ZONE  
& $gcloud compute instances describe $VM --zone $ZONE --format="value(status)"
```

While stopped, you may still be billed for:

- **Persistent disks** (boot disk)
- **Reserved static IPs** (if you created any)

Delete the VM (to remove most costs)

```
& $gcloud --quiet compute instances delete $VM --zone $ZONE
```

Delete the firewall rule \optional cleanup

```
& $gcloud --quiet compute firewall-rules delete $FW_RULE
```

What to tell Cursor next time \copy/paste instructions template

Fill in the [. . .] values.

1) What you want built

- “Build a Python web app with Register + Login. Registration collects: first_name, last_name, address, SSN, username, password. Store in MySQL. On login, fetch user details and show a welcome page.”
- Optional extras (if needed):
 - “Add API /api/add (query params) + UI form/page to call it and show the result.”
 - “Add API /api/employees that calls <https://boringapi.com/api/v1/employees> internally and returns JSON.”
 - “Add a new page employees.html that loads from /api/employees and displays employees in a table. Redirect to it after login.”

2) Tech stack requirements

- “UI must be Python-based (Flask) serving HTML/JS/CSS.”
- “Server must be Python (Flask).”
- “Database must be MySQL.”

3) Containerization requirements

- “Create 3 Docker containers (ui, server, mysql) using Docker Compose.”
- “Expose ports: UI `[8080]`, API `[5000]`, DB `[3306]`.”
- Optional: “Do NOT include phpMyAdmin.” (or “Include phpMyAdmin on `[8081]`.”)

4) Repo + Git instructions

- “Initialize git repo, add ` .gitignore` , commit everything.”
- “Set git identity: `user.name=[...]` , `user.email=[...]` .”
- “Push to remote: ` [https://github.com//.git]` on branch ` main` .”

5) Google Cloud deployment instructions

- “Deploy to GCP Project ID: ` [your-project-id]` .”
- “Stay within free tier / trial credits.”

- “Use Compute Engine `e2-micro` VM + Docker Compose on the VM.”
- “Use zone `[us-central1-a]` (or my preferred zone `[...]`).”
- “Firewall policy: open ports `[8080,5000]` to `[my IP only]` or `[0.0.0.0/0]`.”
- “Repo to deploy on VM: `[your repo URL]`.”
- “Use a startup script to install Docker, create `.env`, and run Docker Compose.”
- “Return the external IP URLs for UI + API health.”
- “Stop the VM when done.”

6) Your local environment details

- “OS: Windows/macOS/Linux. Shell: PowerShell/bash.”
- “Tools installed: Git `[yes/no]` , Google Cloud CLI `[yes/no]` , Docker Desktop `[yes/no]` .”
- “I am logged into gcloud as `[email]` .”
- “Workspace path is `[path]` .”